APEX - Complete User Guide

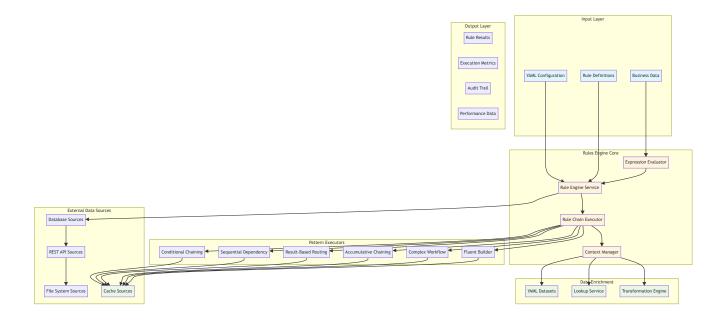
Version: 1.0 Date: 2025-08-02 Author: Mark Andrew Ray-Smith Cityline Ltd

Overview

APEX (Advanced Processing Engine for eXpressions) is a comprehensive rule evaluation system built on expression evaluation technologies with enterprise-grade external data source integration and scenario-based configuration management. It provides a progressive API design that scales from simple rule evaluation to complex business rule management systems with integrated data access capabilities.

Key Capabilities

- Scenario-Based Configuration: Centralized management and routing of data processing pipelines
- YAML Validation System: Enterprise-grade validation with comprehensive error reporting
- External Data Source Integration: Connect to databases, REST APIs, file systems, and caches
- YAML Dataset Enrichment: Embed reference data directly in configuration files
- Progressive API Design: Three-layer API from simple to advanced use cases
- Enterprise Features: Connection pooling, health monitoring, caching, failover
- . High Performance: Optimized for production workloads with comprehensive monitoring



Core Capabilities

Rule Evaluation

- Three-Layer API Design: Simple one-liner evaluation → Structured rule sets → Advanced rule chains
- SpEL Expression Support: Full Spring Expression Language capabilities with custom functions
- Multiple Data Types: Support for primitives, objects, collections, and complex nested structures
- Context Management: Rich evaluation context with variable propagation and result tracking

Configuration Management

- · YAML Configuration: External rule and dataset management with hot-reloading support
- Rule Groups: Organize related rules with execution control and priority management
- Rule Chains: Advanced patterns for nested rules and complex business logic workflows
- Data Service Configuration: Programmatic setup of data sources for rule evaluation
- Metadata Support: Enterprise metadata including business ownership, effective dates, and custom properties

Data Enrichment

- · YAML Dataset Enrichment: Embed lookup datasets directly in configuration files
- Multiple Enrichment Types: Lookup enrichment, transformation enrichment, and custom processors
- Caching Support: High-performance in-memory caching with configurable strategies
- External Integration: Support for database lookups, REST API calls, and custom data sources

Enterprise Features

- Performance Monitoring: Execution time tracking, rule performance analytics, and bottleneck identification
- Error Handling: Comprehensive error management with detailed logging and graceful degradation
- Audit Trail: Complete execution history with rule results and context tracking
- Security: Input validation, expression sandboxing, and access control integration

Financial Services Support

- OTC Derivatives Validation: Specialized rules for financial instrument validation
- Regulatory Compliance: Support for MiFID II, EMIR, and Dodd-Frank requirements
- · Risk Assessment: Multi-criteria risk scoring with weighted components
- Trade Processing: Complex workflow patterns for trade lifecycle management

Advanced Rule Patterns

- Conditional Chaining: Execute expensive rules only when conditions are met
- Sequential Dependency: Build processing pipelines where each stage uses previous results
- · Result-Based Routing: Route to different rule sets based on intermediate results
- Accumulative Chaining: Build up scores across multiple criteria with weighted components
- . Complex Financial Workflow: Multi-stage processing with dependencies and conditional execution
- Fluent Rule Builder: Complex decision trees with conditional branching logic

Architecture Benefits

Developer Experience

- Progressive Complexity: Start simple and add complexity as needed
- Type Safety: Strong typing support with compile-time validation
- Testing Support: Comprehensive testing utilities and mock frameworks

Operations

- Hot Configuration Reload: Update rules without application restart
- · Performance Monitoring: Built-in metrics and monitoring capabilities
- Scalability: Designed for high-throughput, low-latency environments

Business User Friendly

- YAML Configuration: Human-readable configuration format
- Business Metadata: Rich metadata support for business context
- Version Control: Configuration stored in Git with full change history
- Documentation: Self-documenting rules with descriptions and examples

Quick Start (5 Minutes)

Welcome to APEX! This section will get you up and running quickly with three progressively more powerful approaches. Don't worry if some concepts seem unfamiliar at first - we'll explain everything step by step.

Understanding the Basics

Before we dive in, let's understand what APEX does: it evaluates business rules against your data. Think of it as asking questions like "Is this customer old enough?" or "Does this transaction meet our requirements?" APEX uses expressions (written in Spring Expression Language or SpEL) to define these questions.

The # symbol in expressions refers to data you provide. For example, #age >= 18 means "check if the age value is 18 or greater."

1. One-Liner Rule Evaluation (Simplest Approach)

This is the easiest way to get started. You can evaluate a single rule with just one line of code:

```
import dev.mars.rulesengine.core.api.Rules;

// Check if someone is an adult (age 18 or older)
boolean isAdult = Rules.check("#age >= 18", Map.of("age", 25)); // returns true

// Check if account has sufficient balance
boolean hasBalance = Rules.check("#balance > 1000", Map.of("balance", 500)); // returns false

// Working with objects instead of simple values

Customer customer = new Customer("John", 25, "john@example.com");
boolean valid = Rules.check("#data.age >= 18 && #data.email != null", customer); // returns true
```

What's happening here:

- Rules.check() is a static method that evaluates one rule
- The first parameter is your rule expression (the question you're asking)
- The second parameter is your data (either a Map or an object)
- It returns true/false based on whether the rule passes

2. Template-Based Rules (Structured Approach)

When you need multiple related rules, templates provide a cleaner approach:

```
.balanceMinimum(1000)  // Must have at least $1000 balance
.build();

// Validate your customer data against all rules at once
ValidationResult result = validation.validate(customer);

// Check the results
if (result.isValid()) {
    System.out.println("Customer passed all validations!");
} else {
    System.out.println("Validation failed: " + result.getFailureMessages());
}
```

What's happening here:

- RuleSet.validation() creates a builder for common validation scenarios
- Each method (like ageCheck()) adds a pre-configured rule
- build() creates the final rules engine
- validate() runs all rules and gives you a comprehensive result

3. YAML Configuration (Most Flexible Approach)

For complex scenarios or when non-developers need to modify rules, YAML configuration is ideal. This approach separates your business logic from your code and provides the most flexibility.

YAML configuration introduces two powerful concepts that work together:

- Rules: Define your business logic and validation requirements
- Enrichments: Automatically add related data during rule evaluation

Let's explore each concept separately, then see how they work together.

3.1 Rules: Defining Your Business Logic

Rules are the heart of APEX - they define the questions you want to ask about your data. Each rule is like a business requirement written in a way the computer can understand.

Start with a simple rules-only configuration:

```
# Required metadata for all YAML files
metadata:
 name: "Customer Validation Rules"
 version: "1.0.0"
 description: "Basic validation rules for customer data"
 type: "rule-config"
 author: "validation.team@company.com"
# Define your business rules here
rules:
  - id: "age-check"
                                                       # Unique identifier
   name: "Age Validation"
                                                       # Human-readable name
   condition: "#data.age >= 18"
                                                       # The actual rule logic
   message: "Customer must be at least 18 years old" # Error message if rule fails
   severity: "ERROR"
                                                       # How serious is a failure?
  - id: "email-check"
   name: "Email Validation"
   condition: "#data.email != null && #data.email.contains('@')"
```

```
message: "Valid email address is required"
severity: "ERROR"

- id: "name-check"
name: "Name Validation"
condition: "#data.name != null && #data.name.length() > 0"
message: "Customer name is required"
severity: "ERROR"
```

Use this rules configuration in your Java code:

```
// Load the YAML configuration file
RulesEngineConfiguration config = YamlConfigurationLoader.load("customer-rules.yaml");
RulesEngine engine = new RulesEngine(config);
// Prepare your data for evaluation
Map<String, Object> data = Map.of(
   "age", 25,
   "email", "john@example.com",
   "name", "John Doe"
);
// Evaluate all rules
RuleResult result = engine.evaluate(data);
// Check what happened
if (result.isSuccess()) {
   System.out.println("All validation rules passed!");
   System.out.println("Validation failed:");
   result.getFailureMessages().forEach(System.out::println);
}
```

Understanding Rule Conditions:

Rules use Spring Expression Language (SpEL) for conditions:

- #data.age >= 18 Access the 'age' field and check if it's 18 or greater
- #data.email != null Check if the 'email' field exists and is not null
- #data.email.contains('@') Check if the email contains an @ symbol
- #data.name.length() > 0 Check if the name has at least one character

Common Rule Patterns:

```
• Validation: #age >= 18 (must be 18 or older)
```

- Range checking: #score >= 0 && #score <= 100 (score between 0-100)
- Required fields: #email != null (email must exist)
- Pattern matching: #email.contains('@') (email must contain @)
- Complex logic: #amount > 1000 ? #approvalRequired == true : true (amounts over 1000 need approval)

3.2 Enrichments: Adding Smart Data

Enrichments automatically add related information to your data during rule evaluation. Think of them as smart lookups that happen behind the scenes before your rules are evaluated.

Why use enrichments? Instead of just having a status code like "A", enrichments can automatically add the full description "Active Customer" to your data, which your rules can then use.

Create an enrichment-focused configuration:

```
# Required metadata for all YAML files
metadata:
 name: "Customer Data Enrichment"
 version: "1.0.0"
 description: "Enrichment rules to add reference data to customer records"
 type: "rule-config"
 author: "data.enrichment@company.com"
# Enrichments add extra data to your objects during rule evaluation
enrichments:
  - id: "status-enrichment"
                                              # Unique identifier for this enrichment
   name: "Customer Status Enrichment"
                                             # Human-readable name
   type: "lookup-enrichment"
                                              # Type of enrichment (lookup from a dataset)
   condition: "['statusCode'] != null"
                                             # Only enrich if statusCode exists
   lookup-config:
     lookup-dataset:
                                             # The data to look up from
       type: "inline"
                                            # Data is defined right here in the file
       key-field: "code"
                                            # Field to match against
                                            # The actual lookup data
       data:
         - code: "A"
           name: "Active"
           description: "Active customer"
           priority: "High"
          - code: "I"
           name: "Inactive"
           description: "Inactive customer"
           priority: "Low"
          - code: "S"
           name: "Suspended"
           description: "Suspended customer"
           priority: "Medium"
   field-mappings:
                                             # How to add the looked-up data to your object
     - source-field: "name"
                                            # Take the "name" from lookup data
       target-field: "statusName"
                                             # Add it as "statusName" to your object
      - source-field: "description"
       target-field: "statusDescription"
      - source-field: "priority"
       target-field: "customerPriority"
```

Use this enrichment configuration:

```
// Load the enrichment configuration
RulesEngineConfiguration config = YamlConfigurationLoader.load("customer-enrichment.yaml");
RulesEngine engine = new RulesEngine(config);
// Prepare your data (notice we only have the status code)
Map<String, Object> data = Map.of(
    "name", "John Doe",
    "statusCode", "A" // Just the code - enrichment will add more data
);
// Evaluate (this will run enrichments)
RuleResult result = engine.evaluate(data);
// Access the enriched data
Map<String, Object> enrichedData = result.getEnrichedData();
System.out.println("Status Name: " + enrichedData.get("statusName"));
                                                                         // "Active"
System.out.println("Description: " + enrichedData.get("statusDescription")); // "Active customer"
                                                                             // "High"
System.out.println("Priority: " + enrichedData.get("customerPriority"));
```

What happens during enrichment:

- 1. Your data has statusCode: "A"
- 2. APEX looks up "A" in the inline dataset
- 3. It finds the matching record with name "Active", description "Active customer", priority "High"
- 4. It adds statusName: "Active", statusDescription: "Active customer", and customerPriority: "High" to your data
- 5. Your rules can now use these enriched fields

3.3 Combining Rules and Enrichments

The real power of YAML configuration comes from combining rules and enrichments. Enrichments run first to add data, then rules evaluate using both the original and enriched data.

Complete configuration with both rules and enrichments:

```
# Required metadata for all YAML files
metadata:
 name: "Customer Validation and Enrichment"
 version: "1.0.0"
 description: "Complete customer processing with enrichment and validation"
 type: "rule-config"
 author: "customer.processing@company.com"
 created: "2025-08-02"
# Enrichments run FIRST to add reference data
enrichments:
  - id: "status-enrichment"
   name: "Customer Status Enrichment"
   type: "lookup-enrichment"
   condition: "['statusCode'] != null"
   lookup-config:
     lookup-dataset:
       type: "inline"
       key-field: "code"
       data:
          - code: "A"
            name: "Active"
            description: "Active customer"
            allowTransactions: true
            creditLimit: 10000
          - code: "I"
           name: "Inactive"
            description: "Inactive customer"
            allowTransactions: false
            creditLimit: 0
          - code: "S"
            name: "Suspended"
            description: "Suspended customer"
            allowTransactions: false
            creditLimit: 0
   field-mappings:
      - source-field: "name"
       target-field: "statusName"
      - source-field: "allowTransactions"
       target-field: "canTransact"
      - source-field: "creditLimit"
        target-field: "maxCredit"
# Rules run AFTER enrichments and can use the enriched data
rules:
  - id: "age-check"
   name: "Age Validation"
```

```
condition: "#data.age >= 18"
 message: "Customer must be at least 18 years old"
 severity: "ERROR"
- id: "transaction-permission-check"
 name: "Transaction Permission Check"
 condition: "#data.canTransact == true" # Uses enriched field!
 message: "Customer status does not allow transactions"
 severity: "ERROR"
- id: "credit-limit-check"
 name: "Credit Limit Validation"
 condition: "#data.requestedAmount <= #data.maxCredit" # Uses enriched field!</pre>
 message: "Requested amount exceeds customer credit limit"
 severity: "ERROR"
- id: "email-check"
 name: "Email Validation"
 condition: "#data.email != null && #data.email.contains('@')"
 message: "Valid email address is required"
 severity: "WARNING"
```

Using the combined configuration:

```
// Load the complete configuration
RulesEngineConfiguration config = YamlConfigurationLoader.load("customer-complete.yaml");
RulesEngine engine = new RulesEngine(config);
// Prepare your data
Map<String, Object> data = Map.of(
    "age", 25,
    "email", "john@example.com",
    "statusCode", "A", // Will be enriched to add transaction permissions requestedAmount", 5000 // Will be validated against enriched credit limit
);
// Evaluate (enrichments run first, then rules)
RuleResult result = engine.evaluate(data);
// Check results
if (result.isSuccess()) {
    System.out.println("All rules passed!");
    // Access enriched data
    Map<String, Object> enrichedData = result.getEnrichedData();
    System.out.println("Customer Status: " + enrichedData.get("statusName"));
    System.out.println("Can Transact: " + enrichedData.get("canTransact"));
    System.out.println("Credit Limit: " + enrichedData.get("maxCredit"));
} else {
    System.out.println("Some rules failed:");
    result.getFailureMessages().forEach(System.out::println);
}
```

Execution Flow:

1. Enrichment Phase: APEX looks up status code "A" and adds:

```
statusName: "Active"canTransact: truemaxCredit: 10000
```

2. Rules Phase: APEX evaluates rules using both original and enriched data:

- o Age check: 25 >= 18 √ Pass
- o Transaction permission: true == true ✓ Pass (uses enriched canTransact)
- Credit limit: 5000 <= 10000

 √ Pass (uses enriched maxCredit)
- o Email check: john@example.com contains '@' √ Pass

Key Benefits of This Approach:

- Separation of Concerns: Enrichments handle data lookup, rules handle business logic
- . Reusability: The same enrichments can be used by multiple rule sets
- . Maintainability: Business users can modify rules without touching enrichment logic
- Performance: Enrichments run once, rules can use the enriched data multiple times

Which Approach Should You Use?

- One-liner: Perfect for simple, one-off rule checks
- Template-based: Great for common validation scenarios with multiple related rules
- YAML configuration: Best for complex business logic, when rules change frequently, or when business users need to modify
 rules
- YAML with external data sources: Ideal for enterprise scenarios requiring real-time data from databases, APIs, or large datasets

Data Integration Considerations:

- Small, static data: Use inline datasets or external YAML files
- Large or dynamic data: Use external data sources (databases, REST APIs, file systems)
- Enterprise integration: Connect to existing systems using the APEX External Data Sources Guide

You can start with the one-liner approach and gradually move to more sophisticated approaches as your needs grow!

Core Concepts

Now that you've seen the three approaches to using APEX, let's dive deeper into the core concepts that make YAML configuration so powerful. We've already covered Rules and Enrichments in detail in the YAML Configuration section, so here we'll focus on Datasets and provide a quick summary.

Quick Summary: Rules and Enrichments

Rules (detailed in section 3.1):

- Define your business logic using Spring Expression Language (SpEL)
- Include conditions, messages, and severity levels
- Run after enrichments and can use enriched data
- Common patterns: validation, range checking, required fields

Enrichments (detailed in section 3.2):

- · Automatically add related data during rule evaluation
- · Run before rules to provide additional context
- · Use lookup datasets to find related information
- Map lookup results to new fields in your data

Combined Power (detailed in section 3.3):

- · Enrichments run first to add reference data
- · Rules then evaluate using both original and enriched data
- · Enables sophisticated business logic with clean separation of concerns

Datasets: Your Reference Data

Datasets are collections of reference data that enrichments use for lookups. They're like lookup tables that contain additional information about codes, IDs, or other identifiers in your data.

Two ways to organize your datasets:

Inline Datasets (Best for small, unique data)

Use inline datasets when you have small amounts of data that are specific to one configuration file:

When to use inline datasets:

- Small datasets (less than 50 records)
- · Data that's unique to this specific configuration
- · Simple lookup tables that won't be reused elsewhere

External Dataset Files (Best for larger, reusable data)

Use external files when you have larger datasets or data that multiple configurations might use:

Then create datasets/currencies.yaml:

```
# Required metadata for all YAML files
metadata:
    name: "Currency Reference Data"
    version: "1.0.0"
    description: "Currency codes with names, symbols, and regional information"
    type: "dataset"
    source: "ISO 4217 currency codes"
    last-updated: "2025-08-02"
```

```
# The actual dataset
  - code: "USD"
   name: "US Dollar"
   symbol: "$"
   region: "North America"
  - code: "EUR"
    name: "Euro"
   symbol: "€"
   region: "Europe"
  - code: "GBP"
   name: "British Pound"
    symbol: "£"
   region: "Europe"
  - code: "JPY"
    name: "Japanese Yen"
    symbol: "\forall"
    region: "Asia"
```

When to use external dataset files:

- · Larger datasets (50+ records)
- · Data that multiple configurations need to share
- · Data that changes independently of your rule configurations
- · When you want to keep your main configuration file clean and focused

YAML Configuration Guide

YAML (Yet Another Markup Language) is a human-readable format for configuration files. Don't worry if you're new to YAML - it's designed to be easy to read and write. Think of it as a structured way to organize information, similar to how you might organize information in an outline.

Understanding YAML Basics

YAML uses indentation (spaces) to show relationships between items. Here are the key concepts:

- Indentation matters: Use spaces (not tabs) to show hierarchy
- Lists: Items that start with a dash ()
- Key-value pairs: key: value
- · Nested structures: Indent to show items belong together

Configuration Structure

Every APEX configuration file follows this basic structure:

```
# Metadata section: Information about this configuration file
metadata:
    name: "Configuration Name"  # What this configuration does
    version: "1.0.0"  # Version for tracking changes
    description: "Configuration description" # Detailed explanation
    type: "rule-config"  # Required: File type (rule-config, scenario, dataset, etc.)
    author: "Team Name"  # Who created/maintains this
    created: "2025-08-02"  # When it was created
    last-modified: "2025-08-02"  # Last update date
```

```
tags: ["tag1", "tag2"]  # Categories for organization

# Rules section: Your business logic
rules:
    # Individual rule definitions go here
    # Each rule defines a condition to check

# Enrichments section: Data enhancement
enrichments:
    # Enrichment definitions go here
    # Each enrichment adds data to your objects

# Rule groups section: Organized rule collections
rule-groups:
    # Rule group definitions go here
    # Groups let you organize related rules together
```

Why organize it this way?

- Metadata: Helps you track and document your configurations
- Rules: Contains your business logic and validation requirements
- . Enrichments: Automatically adds useful information to your data
- · Rule Groups: Organizes related rules for better management

Rule Configuration

Rules are where you define your business logic. Each rule is like a question you're asking about your data. Here's how to configure them:

```
rules:
  - id: "unique-rule-id"
                                             # Required: Unique identifier (like a name tag)
                                         # Required: What this rule does in plain English
   name: "Human Readable Name"
   condition: "#data.field > 100"
message: "Validation message"
                                         # Required: The actual business logic to check
# Optional: What to show if the rule fails
   severity: "ERROR"
                                             # Optional: How serious is a failure? (ERROR, WARNING, INFO)
   enabled: true
                                             # Optional: Turn this rule on/off (default: true)
   tags: ["validation", "business"] # Optional: Categories for organization
                                             # Optional: Additional information for governance
   metadata:
     owner: "Business Team"
                                          # Who owns/maintains this rule
      domain: "Finance"
                                          # What business area it belongs to
      purpose: "Regulatory compliance" # Why this rule exists
```

Understanding each part:

- id: A unique name for this rule (like "customer-age-check"). Use descriptive names that make sense to your team.
- name: A human-friendly description that anyone can understand (like "Customer Age Validation").
- condition: The actual business logic using SpEL expressions. Common patterns:
 - #age >= 18 (age must be 18 or older)
 - o #amount > 0 && #amount <= 1000 (amount between 0 and 1000)
 - #email != null && #email.contains('@') (email must exist and contain @)
- message: What users see when the rule fails. Make it helpful and actionable.
- severity: How important is this rule?
 - o ERROR: Critical must be fixed
 - WARNING: Important should be reviewed
 - INFO: Informational good to know

- enabled: Allows you to temporarily turn rules on/off without deleting them.
- tags: Help organize and filter rules. Use consistent tags across your organization.
- metadata: Additional information for governance, documentation, and audit trails.

Enrichment Configuration

Enrichments automatically add related information to your data. Think of them as smart lookups that happen automatically during rule evaluation.

```
enrichments:
   - id: "enrichment-id"
                                                                     # Required: Unique identifier for this enrichment
     id: "enrichment-id"  # Required: Unique identifier for this enrichment
type: "lookup-enrichment"  # Required: Type of enrichment (lookup is most common)
condition: "['field'] != null"  # Optional: Only enrich if this condition is true
enabled: true  # Optional: Turn this enrichment on/off (default: true)
      enabled: true
                                                                      # Optional: Turn this enrichment on/off (default: true)
           kup-config: # Configuration for the lookup process

ookup-dataset: # Where to find the lookup data

type: "inline" # Data source type: "inline" or "yaml-file"

key-field: "lookupKey" # Field to match against in your data

cache-enabled: true # Keep lookup data in memory for speed

cache-ttl-seconds: 3600 # How long to cache (1 hour = 3600 seconds)

default-values: # What to use if no match is found
      lookup-config:
        lookup-dataset:
               defaultField: "defaultValue"
               ta:
- lookupKey: "key1"
- "value1"
                                                                 # The actual lookup data (for inline type)
            data:
                                                                   # This is what we match against
                                                                   # Additional data to add
                   field2: "value2"
      field-mappings:
                                                                 # How to add the looked-up data to your object
         - source-field: "field1"
                                                                 # Take this field from the lookup data
           target-field: "enrichedField1"  # Add it to your object with this name
         - source-field: "field2"
            target-field: "enrichedField2"
```

Understanding enrichment flow:

- 1. Check condition: If specified, only enrich when the condition is true
- 2. Find matching data: Look up the key value in your dataset
- 3. Map fields: Copy specified fields from the lookup data to your object
- 4. Use defaults: If no match found, use default values (if configured)

Example in action:

```
    Your data has: {statusCode: "A"}
    Lookup dataset has: {code: "A", name: "Active", description: "Customer is active"}
    Field mappings copy name to statusName and description to statusDescription
    Result: Your data now has: {statusCode: "A", statusName: "Active", statusDescription: "Customer is active"}
```

Common use cases:

- Convert codes to human-readable names (status codes, country codes, etc.)
- Add regional information based on location codes
- · Enrich product data with category information
- · Add calculated fields based on lookup tables

Dataset Enrichment

Dataset enrichment is one of APEX's most powerful features. It automatically adds related information to your data during rule evaluation, transforming simple codes into rich, meaningful data.

Understanding Dataset Enrichment

Imagine you have customer data with just a status code like "A". Dataset enrichment can automatically add the full status name "Active" and description "Customer account is active and in good standing" to your data. This happens transparently during rule evaluation, so your rules can work with both the original code and the enriched information.

When to Use Dataset Enrichment

Dataset enrichment works best for reference data - information that helps explain or categorize your main data.

Perfect candidates for dataset enrichment:

- Currency codes and names: "USD" → "US Dollar", "EUR" → "Euro"
- Country codes and regions: "US" → "United States", "North America"
- Status codes and descriptions: "A" → "Active", "Customer account is active"
- Product categories: "ELEC" → "Electronics", "Consumer electronics category"
- · Reference data that changes infrequently: Data that's stable over time
- Small to medium datasets: Less than 1000 records for optimal performance

Not suitable for dataset enrichment:

- Large datasets: More than 1000 records (use external data sources instead)
- . Frequently changing data: Data that updates multiple times per day
- Data requiring complex business logic: Calculations or complex transformations
- · Real-time data from external systems: Live data that needs fresh API calls

Why these limitations? Dataset enrichment loads all data into memory for fast lookups. This works great for small, stable reference data but isn't efficient for large or frequently changing datasets.

Solution for Large or Dynamic Data: For scenarios that exceed these limitations, APEX provides powerful external data sources that can connect to databases, REST APIs, file systems, and caches. See the complete APEX External Data Sources Guide for enterprise-scale data integration.

Dataset Types

1. Inline Datasets

Best for small, unique datasets:

```
lookup-dataset:
  type: "inline"
  key-field: "code"
  data:
    - code: "A"
     name: "Active"
    - code: "I"
     name: "Inactive"
```

2. External YAML Files

Best for reusable datasets:

Create datasets/statuses.yaml:

```
data:
    - code: "A"
    name: "Active"
    description: "Active status"
    code: "I"
    name: "Inactive"
    description: "Inactive status"
```

Reference in configuration:

```
lookup-dataset:
  type: "yaml-file"
  file-path: "datasets/statuses.yaml"
  key-field: "code"
```

3. External Data Sources

For dynamic, large-scale, or real-time data, APEX provides powerful external data source integration. Instead of storing data in YAML files, you can connect directly to databases, REST APIs, file systems, and caches.

Why use external data sources?

- · Live data: Always get the most current information
- · Large datasets: Handle millions of records efficiently
- Real-time updates: Data changes are immediately available
- Enterprise integration: Connect to existing systems and databases
- Performance: Optimized queries and caching for high-volume operations

Database Integration Example:

```
# Configuration with database lookup
metadata:
  name: "Customer Database Enrichment"
  version: "1.0.0"
  description: "Customer data enrichment using PostgreSQL database"
  type: "rule-config"
  author: "data.integration@company.com"
enrichments:
  - id: "customer-enrichment"
    name: "Customer Database Lookup"
    type: "lookup-enrichment"
    condition: "['customerId'] != null"
    lookup-config:
      lookup-dataset:
        type: "database"
                                            # External database source
        connection-name: "customer-db"
                                           # Named connection
        query: "SELECT customer_name, credit_rating, account_status FROM customers WHERE customer_id = ?"
       key-field: "customerId"  # Field to use in WHERE clause
cache-enabled: true  # Cache results for performance
        cache-enabled: true
                                         # Cache results for performance
        cache-ttl-seconds: 300  # Cache for 5 minutes
    field-mappings:
```

```
    source-field: "customer_name" target-field: "customerName"
    source-field: "credit_rating" target-field: "creditRating"
    source-field: "account_status" target-field: "accountStatus"
```

REST API Integration Example:

```
# Configuration with REST API lookup
enrichments:
 - id: "market-data-enrichment"
   name: "Real-time Market Data"
   type: "lookup-enrichment"
   condition: "['symbol'] != null"
   lookup-config:
     lookup-dataset:
       type: "rest-api"
                                          # External REST API source
       base-url: "https://api.marketdata.com"
       endpoint: "/v1/quotes/{symbol}"  # {symbol} will be replaced
       method: "GET"
       headers:
         Authorization: "Bearer ${API_TOKEN}" # Environment variable
         Content-Type: "application/json"
       timeout-seconds: 5
       cache-enabled: true
        cache-ttl-seconds: 60
                                  # Cache for 1 minute
   field-mappings:
      - source-field: "price"
       target-field: "currentPrice"
      - source-field: "volume"
       target-field: "tradingVolume"
```

File System Integration Example:

```
# Configuration with file system lookup
enrichments:
  - id: "reference-data-enrichment"
   name: "File System Reference Data"
   type: "lookup-enrichment"
   condition: "['productCode'] != null"
   lookup-config:
     lookup-dataset:
       type: "file-system"
                                         # External file system source
       file-path: "/data/products/products.csv"
       format: "csv"
                                        # CSV, JSON, XML supported
       key-field: "code"
       delimiter: ","
       has-header: true
       cache-enabled: true
       watch-for-changes: true
                                   # Reload when file changes
   field-mappings:
     - source-field: "name"
       target-field: "productName"
      - source-field: "category"
       target-field: "productCategory"
```

Using External Data Sources in Java:

```
// Configure external data sources
ExternalDataSourceConfiguration dbConfig = ExternalDataSourceConfiguration.builder()
    .connectionName("customer-db")
    .jdbcUrl("jdbc:postgresql://localhost:5432/customers")
    .username("${DB_USERNAME}") // From environment variables
    .password("${DB_PASSWORD}")
    .build();
// Load configuration with external data sources
RulesEngineConfiguration config = YamlConfigurationLoader.load("customer-enrichment.yaml");
RulesEngine engine = new RulesEngine(config);
// Register external data sources
engine.registerDataSource("customer-db", dbConfig);
// Use normally - external data sources work transparently
Map<String, Object> data = Map.of("customerId", "CUST123");
RuleResult result = engine.evaluate(data);
// Access enriched data from database
String customerName = (String) result.getEnrichedData().get("customerName");
String creditRating = (String) result.getEnrichedData().get("creditRating");
```

Key Benefits of External Data Sources:

- . Always Current: Data is fetched in real-time or from cache
- Scalable: Handle large datasets without memory constraints
- Enterprise Ready: Connect to existing databases and APIs
- · Performance Optimized: Built-in caching and connection pooling
- Secure: Support for authentication and encrypted connections

When to Use Each Type:

Dataset Type	Best For	Example Use Cases		
Inline	Inline Small, static data Status codes, country codes, simple loc			
External YAML	Medium, shared data	Product catalogs, reference data, configuration		
Database	Large, dynamic data	Customer records, transaction history, real-time data		
REST API	External services	Market data, third-party services, microservices		
File System	Batch data, reports	Daily files, CSV imports, data feeds		

■ For Complete External Data Sources Documentation:

This is just an introduction to external data sources. For comprehensive coverage including:

- · Database connection configuration and pooling
- REST API authentication and error handling
- · File system monitoring and format support
- · Caching strategies and performance optimization
- Security and encryption
- · Enterprise patterns and best practices

Performance Optimization

```
lookup-dataset:
  type: "inline"
  key-field: "code"
  cache-enabled: true
  cache-ttl-seconds: 3600
  preload-enabled: true
  data:
    # Dataset entries
```

Scenario-Based Configuration Management

Overview

APEX includes a powerful scenario-based configuration system that provides centralized management and routing of data processing pipelines. This system enables organizations to manage complex rule configurations through lightweight scenario files that associate data types with appropriate rule configurations.

Key Benefits

- Centralized Management: Single registry manages all available scenarios
- Lightweight Configuration: Scenario files contain only routing information, not business logic
- Type-Safe Processing: Automatic routing based on data type detection
- Dependency Tracking: Complete analysis of configuration file dependencies
- Enterprise Validation: Comprehensive YAML validation with detailed error reporting

How Scenarios Work

Scenarios provide a three-layer architecture:

- 1. Registry Layer: Central discovery mechanism (config/data-type-scenarios.yaml)
- 2. Scenario Layer: Lightweight routing files (scenarios/*.yaml)
- 3. Rule Configuration Layer: Actual business logic files (bootstrap/*.yaml , config/*.yaml)

Example: OTC Options Scenario

Here's how to set up a scenario for processing OTC Options:

1. Registry Entry (config/data-type-scenarios.yaml):

```
scenario-registry:
    scenario-id: "otc-options-standard"
    config-file: "scenarios/otc-options-scenario.yaml"
    data-types: ["OtcOption", "dev.mars.apex.demo.bootstrap.model.OtcOption"]
    description: "Standard validation and enrichment pipeline for OTC Options"
    business-domain: "Derivatives Trading"
    owner: "derivatives.team@company.com"
```

2. Scenario File (scenarios/otc-options-scenario.yaml):

```
metadata:
 name: "OTC Options Processing Scenario"
 version: "1.0.0"
 description: "Associates OTC Options with existing rule configurations"
 type: "scenario"
 business-domain: "Derivatives Trading"
 owner: "derivatives.team@company.com"
scenario:
 scenario-id: "otc-options-standard"
 name: "OTC Options Standard Processing"
 description: "Associates OTC Options with existing rule configurations"
 # Data types this scenario applies to
 data-types:
    - "dev.mars.apex.demo.bootstrap.model.OtcOption"
   - "OtcOption" # Short name alias
 # References to existing rule configuration files
 rule-configurations:
    - "bootstrap/otc-options-bootstrap.yaml"
   - "config/derivatives-validation-rules.yaml"
```

3. Using Scenarios in Code:

```
// Load the scenario registry
DataTypeScenarioService scenarioService = new DataTypeScenarioService();
scenarioService.loadScenarios("config/data-type-scenarios.yaml");

// Automatic routing based on data type
OtcOption option = new OtcOption(...);
ScenarioConfiguration scenario = scenarioService.getScenarioForData(option);

// Get rule configuration files to load and execute
List<String> ruleFiles = scenario.getRuleConfigurations();
for (String ruleFile : ruleFiles) {
    RuleConfiguration rules = ruleLoader.load(ruleFile);
    ruleEngine.execute(rules, option);
}
```

YAML Validation System

APEX includes comprehensive YAML validation to ensure configuration integrity:

Validation Features:

- Required Metadata: All YAML files must include proper metadata with required type field
- Type-Specific Validation: Different validation rules for different file types
- Dependency Validation: Validates complete dependency chains and detects missing references
- Syntax Validation: Ensures proper YAML syntax and structure
- · Comprehensive Reporting: Detailed validation reports with errors, warnings, and recommendations

Mandatory Metadata Attributes:

All YAML files in APEX must include a metadata section with the following required fields. For complete standards and examples, see the Configuration Standards and Validation section.

```
metadata:
  name: "Descriptive Name"  # Required: Human-readable name
  version: "1.0.0"  # Required: Semantic version number
  description: "Clear description"  # Required: What this file does
  type: "file-type"  # Required: One of the supported types below
```

Additional Required Fields by Type:

File Type	Additional Required Fields	Purpose
scenario	business-domain , owner	Business context and ownership
scenario-registry	created-by	Registry management
bootstrap	business-domain , created-by	Demo context and creator
rule-config	author	Rule authorship
dataset	source	Data source information
enrichment	author	Enrichment logic authorship
rule-chain	author	Rule chain authorship

Supported File Types:

scenario: Scenario configuration files

scenario-registry : Central scenario registry

• bootstrap : Complete demo configurations

• rule-config: Reusable rule configurations

• dataset : Data reference files

enrichment : Data enrichment configurations

rule-chain: Sequential rule execution files

Example Complete Metadata:

```
# Example for a scenario file
metadata:
    name: "OTC Options Processing Scenario"
    version: "1.0.0"
    description: "Associates OTC Options with existing rule configurations"
    type: "scenario"
    business-domain: "Derivatives Trading"
    owner: "derivatives.team@company.com"
    created: "2025-08-02"
    last-modified: "2025-08-02"
    tags: ["derivatives", "options", "trading"]

# Example for a rule-config file
metadata:
    name: "Financial Validation Rules"
    version: "1.0.0"
    description: "Comprehensive validation rules for financial instruments"
```

```
type: "rule-config"
author: "rules.team@company.com"
created: "2025-08-02"
business-domain: "Financial Services"
```

Example Validation Usage:

```
// Validate a single file
YamlMetadataValidator validator = new YamlMetadataValidator();
YamlValidationResult result = validator.validateFile("scenarios/otc-options-scenario.yaml");

if (result.isValid()) {
    System.out.println(" / File is valid");
} else {
    System.out.println(" / Validation errors:");
    result.getErrors().forEach(error -> System.out.println(" - " + error));
}

// Validate multiple files
List<String> files = List.of("scenarios/scenario1.yaml", "config/rules.yaml");
YamlValidationSummary summary = validator.validateFiles(files);
String report = summary.getReport(); // Comprehensive validation report
```

Available Scenarios

APEX includes several pre-built scenarios for common financial services use cases:

- OTC Options Standard Processing: Complete validation and enrichment pipeline for OTC Options
- Commodity Swaps Standard Processing: Multi-layered validation for commodity derivatives
- Settlement Auto-Repair: Intelligent auto-repair for failed settlement instructions

Best Practices

- 1. Keep Scenarios Lightweight: Only include data type mappings and rule file references
- 2. Use Descriptive Names: Clear scenario IDs and descriptions for maintainability
- 3. Validate Regularly: Run YAML validation as part of CI/CD pipeline
- 4. Document Ownership: Include clear ownership and contact information
- 5. Version Control: Use semantic versioning for scenario configurations

For comprehensive configuration standards, naming conventions, and implementation checklists, see the Configuration Standards and Validation section.

External Data Source Integration

Overview

The APEX Rules Engine provides comprehensive external data source integration, enabling seamless access to databases, REST APIs, file systems, and caches through a unified interface. This enterprise-grade integration supports advanced features like connection pooling, health monitoring, caching, and automatic failover.

Supported Data Source Types

1. Database Sources

Connect to relational databases with full connection pooling support:

```
dataSources:
 - name: "user-database"
   type: "database"
   sourceType: "postgresql"
   enabled: true
   connection:
     host: "localhost"
     port: 5432
     database: "myapp"
     username: "app_user"
     password: "${DB_PASSWORD}"
     maxPoolSize: 20
     minPoolSize: 5
   queries:
     getUserById: "SELECT * FROM users WHERE id = :id"
     getAllUsers: "SELECT * FROM users ORDER BY created_at DESC"
   parameterNames:
      - "id"
   cache:
     enabled: true
     ttlSeconds: 300
     maxSize: 1000
```

Supported Databases: PostgreSQL, MySQL, Oracle, SQL Server, H2

2. REST API Sources

Integrate with HTTP/HTTPS APIs with various authentication methods:

```
dataSources:
 - name: "external-api"
   type: "rest-api"
   enabled: true
   connection:
     baseUrl: "https://api.example.com/v1"
     timeout: 10000
     retryAttempts: 3
   authentication:
     type: "bearer"
     token: "${API_TOKEN}"
   endpoints:
     getUser: "/users/{userId}"
     searchUsers: "/users/search?q={query}"
   parameterNames:
      - "userId"
     - "query"
   circuitBreaker:
     enabled: true
```

```
failureThreshold: 5
recoveryTimeout: 30000
```

Authentication Types: Bearer tokens, API keys, Basic auth, OAuth2

3. File System Sources

Process various file formats with automatic parsing:

```
dataSources:
  - name: "data-files"
    type: "file-system"
    enabled: true
    connection:
     basePath: "/data/files"
     filePattern: "*.csv"
      watchForChanges: true
      encoding: "UTF-8"
    fileFormat:
      type: "csv"
      hasHeaderRow: true
      delimiter: ","
      columnMappings:
        "customer_id": "id"
        "customer_name": "name"
    parameterNames:
      - "filename"
```

Supported Formats: CSV, JSON, XML, fixed-width, plain text

4. Cache Sources

High-performance in-memory caching:

```
dataSources:
    name: "app-cache"
    type: "cache"
    sourceType: "memory"
    enabled: true

    cache:
        enabled: true
    maxSize: 10000
    ttlSeconds: 1800
    evictionPolicy: "LRU"
```

Using External Data Sources

Basic Usage

```
// Initialize configuration service
DataSourceConfigurationService configService = DataSourceConfigurationService.getInstance();
```

```
YamlRuleConfiguration yamlConfig = loadConfiguration("data-sources.yaml");
configService.initialize(yamlConfig);

// Get data source
ExternalDataSource userDb = configService.getDataSource("user-database");

// Execute queries
Map<String, Object> parameters = Map.of("id", 123);
List<Object> results = userDb.query("getUserById", parameters);

// Get single result
Object user = userDb.queryForObject("getUserById", parameters);
```

Advanced Usage with Load Balancing

```
// Get manager for advanced operations
DataSourceManager manager = configService.getDataSourceManager();

// Load balancing across multiple sources
ExternalDataSource source = manager.getDataSourceWithLoadBalancing(DataSourceType.DATABASE);

// Failover query across healthy sources
List<Object> results = manager.queryWithFailover(DataSourceType.DATABASE, "getAllUsers", Collections.emptyMap());

// Async operations
CompletableFuture<List<Object>> future = manager.queryAsync("user-database", "getAllUsers", Collections.emptyMap());
List<Object> users = future.get(10, TimeUnit.SECONDS);
```

Enterprise Features

Health Monitoring

```
healthCheck:
   enabled: true
   intervalSeconds: 30
   timeoutSeconds: 5
   failureThreshold: 3
   query: "SELECT 1"
```

Environment-Specific Configuration

Monitoring and Statistics

```
// Get performance metrics
DataSourceMetrics metrics = dataSource.getMetrics();
System.out.println("Success rate: " + metrics.getSuccessRate());
System.out.println("Average response time: " + metrics.getAverageResponseTime());
// Registry statistics
RegistryStatistics stats = registry.getStatistics();
System.out.println("Health percentage: " + stats.getHealthPercentage());
```

Integration with Rules

External data sources integrate seamlessly with the rules engine:

```
# Use data sources in rule conditions
rules:
 - id: "user-validation"
   condition: "dataSource('user-database').queryForObject('getUserById', {'id': #userId}) != null"
   message: "User exists in database"
# Use in enrichments
enrichments:
 - id: "user-enrichment"
   type: "data-source-enrichment"
   data-source: "user-database"
   query: "getUserById"
   parameters:
     id: "#userId"
   field-mappings:
     - source-field: "name"
       target-field: "userName"
     - source-field: "email"
       target-field: "userEmail"
```

Best Practices

Configuration Management

- Use environment variables for sensitive data
- Implement environment-specific overrides
- · Validate configurations before deployment
- · Use meaningful, descriptive names

Performance Optimization

- · Configure appropriate connection pool sizes
- · Enable caching for frequently accessed data
- · Use circuit breakers for external APIs
- · Monitor performance metrics regularly

Security

- Always use SSL/TLS in production
- · Implement proper access controls

- · Use strong authentication methods
- Encrypt sensitive configuration data

Error Handling

- · Configure health checks appropriately
- · Implement retry logic with exponential backoff
- Use graceful degradation strategies
- · Monitor and alert on failures

For detailed configuration guides, see:

- Database Configuration Guide
- REST API Configuration Guide
- File System Configuration Guide
- · Best Practices Guide

Rule Groups Configuration

Overview

Rule Groups provide a way to organize related rules and control their execution as a logical unit. Rules within a group can be combined using AND or OR operators, allowing for complex validation scenarios where multiple conditions must be met (AND) or where any one of several conditions is sufficient (OR).

Key Features

- Logical Operators: Combine rules with AND or OR operators
- Priority Management: Control execution order within groups
- · Category Support: Organize groups by business domain
- Sequence Control: Define rule execution order within groups
- Metadata Support: Rich metadata for governance and audit trails

Programmatic Rule Group Creation

Using RuleGroupBuilder

```
// Create a rule group with AND operator
RuleGroup validationGroup = new RuleGroupBuilder()
    .withId("validation-group")
    .withName("Customer Validation Group")
    .withDescription("Complete customer validation checks")
    .withCategory("customer-validation")
    .withPriority(10)
    .withAndOperator() // All rules must pass
    .build();
// Create a rule group with OR operator
RuleGroup eligibilityGroup = new RuleGroupBuilder()
    .withId("eligibility-group")
    .withName("Customer Eligibility Group")
    .withDescription("Customer eligibility checks")
    .withCategory("customer-eligibility")
    .withPriority(20)
```

```
.withOrOperator() // Any rule can pass
.build();
```

Using RulesEngineConfiguration

```
RulesEngineConfiguration config = new RulesEngineConfiguration();
// Create rule group with AND operator
RuleGroup andGroup = config.createRuleGroupWithAnd(
   "RG001",
                                      // Group ID
   new Category("validation", 10),
                                      // Category
   "Validation Checks",
                                      // Name
   "All validation rules must pass", // Description
                                       // Priority
);
// Create rule group with OR operator
RuleGroup orGroup = config.createRuleGroupWithOr(
   "RG002",
                                      // Group ID
   new Category("eligibility", 20), // Category
   "Eligibility Checks",
                                      // Name
   "Any eligibility rule can pass", // Description
                                      // Priority
);
// Create multi-category rule group
Set<String> categories = Set.of("validation", "compliance");
RuleGroup multiCategoryGroup = config.createRuleGroupWithAnd(
   "RG003",
                                      // Group ID
   categories,
                                      // Multiple categories
   "Compliance Validation",
                                      // Name
   "Compliance and validation checks", // Description
                                      // Priority
);
```

Adding Rules to Groups

```
// Create individual rules
Rule ageRule = config.rule("age-check")
    .withName("Age Validation")
    .withCondition("#age >= 18")
    .withMessage("Customer must be at least 18")
    .build();
Rule emailRule = config.rule("email-check")
    .withName("Email Validation")
    .withCondition("#email != null && #email.contains('@')")
    .withMessage("Valid email required")
    .build();
Rule incomeRule = config.rule("income-check")
    .withName("Income Validation")
    .withCondition("#income >= 25000")
    .withMessage("Minimum income requirement")
    .build();
// Add rules to AND group (all must pass)
andGroup.addRule(ageRule, 1); // Execute first
andGroup.addRule(emailRule, 2); // Execute second
andGroup.addRule(incomeRule, 3); // Execute third
```

```
// Add rules to OR group (any can pass)
orGroup.addRule(ageRule, 1);
orGroup.addRule(incomeRule, 2);

// Register groups with configuration
config.registerRuleGroup(andGroup);
config.registerRuleGroup(orGroup);
```

YAML Rule Group Configuration

Basic Rule Group Configuration

```
metadata:
 name: "Customer Processing Rules"
 version: "1.0.0"
 description: "Customer validation and processing rules"
 type: "rule-config"
 author: "customer.team@company.com"
rules:
  - id: "age-validation"
   name: "Age Check"
   condition: "#age >= 18"
   message: "Customer must be at least 18"
  - id: "email-validation"
   name: "Email Check"
   condition: "#email != null && #email.contains('@')"
   message: "Valid email address required"
  - id: "income-validation"
   name: "Income Check"
   condition: "#income >= 25000"
   message: "Minimum income of $25,000 required"
rule-groups:
  - id: "customer-validation"
   name: "Customer Validation Rules"
   description: "Complete customer validation rule set"
   category: "validation"
   priority: 10
   enabled: true
   stop-on-first-failure: false
   parallel-execution: false
   rule-ids:
     - "age-validation"
     - "email-validation"
     - "income-validation"
   metadata:
     owner: "Customer Team"
     domain: "Customer Management"
     purpose: "Customer data validation"
```

Advanced Rule Group Configuration

```
rule-groups:
    # AND group - all rules must pass
- id: "strict-validation"
    name: "Strict Validation Group"
    description: "All validation rules must pass"
```

```
category: "validation"
 categories: ["validation", "compliance"] # Multiple categories
 enabled: true
 stop-on-first-failure: true # Stop on first failure for efficiency
 parallel-execution: false  # Sequential execution
   - "age-validation"
   - "email-validation"
   - "income-validation"
 tags: ["strict", "validation", "required"]
 metadata:
   owner: "Compliance Team"
   business-domain: "Customer Onboarding"
   created-by: "compliance.admin@company.com"
 execution-config:
   timeout-ms: 5000
   retry-count: 3
   circuit-breaker: true
# OR group - any rule can pass with advanced rule references
- id: "eligibility-check"
 name: "Customer Eligibility Check"
 description: "Customer meets at least one eligibility criteria"
 category: "eligibility"
 priority: 20
 enabled: true
 stop-on-first-failure: false # Continue even if one rule fails
 parallel-execution: true  # Parallel execution for performance
 rule-references:
   - rule-id: "premium-customer"
     sequence: 1
     enabled: true
     override-priority: 5
    - rule-id: "long-term-customer"
     sequence: 2
     enabled: true
     override-priority: 10
    - rule-id: "high-value-customer"
     sequence: 3
      enabled: true
     override-priority: 15
 metadata:
   owner: "Business Team"
   purpose: "Customer eligibility determination"
```

Complete YAML Configuration Reference

```
rule-groups:
  - id: "complete-example"
                                             # Required: Unique identifier
   name: "Complete Rule Group Example" # Required: Human-readable name
   description: "Shows all available options" # Optional: Description
   # Category Configuration
   category: "validation"
                                             # Single category
   categories: ["validation", "compliance"] # Multiple categories (alternative to category)
   # Execution Configuration
   priority: 10
                                            # Optional: Execution priority (default: 100)
   enabled: true
                                           # Optional: Enable/disable group (default: true)
   stop-on-first-failure: false
                                           # Optional: Stop on first rule failure (default: false)
   parallel-execution: false
                                           # Optional: Execute rules in parallel (default: false)
```

```
# Rule References - Option 1: Simple rule IDs
rule-ids:
 - "rule-1"
 - "rule-2"
 - "rule-3"
# Rule References - Option 2: Advanced rule references with control
rule-references:
  - rule-id: "advanced-rule-1"
   sequence: 1
                                        # Optional: Execution sequence
   enabled: true
                                        # Optional: Enable/disable this rule (default: true)
   override-priority: 5
                                        # Optional: Override rule's default priority
  - rule-id: "advanced-rule-2"
   sequence: 2
    enabled: false
                                        # Disabled rule
   override-priority: 10
# Metadata and Tags
tags: ["validation", "customer", "strict"] # Optional: Tags for categorization
metadata:
                                        # Optional: Custom metadata
 owner: "Team Name"
 business-domain: "Domain"
 created-by: "user@company.com"
 purpose: "Business purpose"
 custom-field: "custom-value"
# Advanced Execution Configuration
execution-config:
                                       # Optional: Advanced execution settings
                                       # Optional: Timeout in milliseconds
 timeout-ms: 5000
                                      # Optional: Number of retries on failure
 retry-count: 3
 circuit-breaker: true
                                       # Optional: Enable circuit breaker pattern
```

YAML Configuration Properties Reference

Property	Туре	Required	Default	Description
id	String	Yes	-	Unique identifier for the rule group
name	String	Yes	-	Human-readable name for the rule group
description	String	No	""	Description of what the rule group does
category	String	No	-	Single category for the rule group
categories	List	No	-	Multiple categories (alternative to category)
priority	Integer	No	100	Execution priority (lower = higher priority)
enabled	Boolean	No	true	Whether the rule group is enabled
stop-on-first- failure	Boolean	No	false	Stop execution on first rule failure (AND logic)
parallel-execution	Boolean	No	false	Execute rules in parallel for performance
rule-ids	List	No	-	Simple list of rule IDs to include
rule-references	List	No	-	Advanced rule references with control options

Property	Туре	Required	Default	Description
tags	List	No	-	Tags for categorization and filtering
metadata	Map <string, object=""></string,>	No	-	Custom metadata for governance
execution-config	ExecutionConfig	No	-	Advanced execution configuration

Rule Reference Properties

Property	Туре	Required	Default	Description
rule-id	String	Yes	-	ID of the rule to reference
sequence	Integer	No	-	Execution sequence within the group
enabled	Boolean	No	true	Whether this rule is enabled in the group
override-priority	Integer	No	-	Override the rule's default priority

Execution Config Properties

Property	Туре	Required	Default	Description
timeout-ms	Long	No	-	Timeout in milliseconds for rule group execution
retry-count	Integer	No	-	Number of retries on execution failure
circuit-breaker	Boolean	No	false	Enable circuit breaker pattern for resilience

Rule Group Execution Behavior

AND Groups (All Rules Must Pass)

```
// AND group behavior
RuleGroup andGroup = config.createRuleGroupWithAnd("and-group", category, "AND Group", "All must pass", 10);

// Execution logic:
// - All rules must evaluate to true
// - If any rule fails, the entire group fails
// - Execution can stop on first failure (configurable)
// - Result is true only if ALL rules pass

Map<String, Object> data = Map.of("age", 25, "email", "test@example.com", "income", 30000);
RuleResult result = engine.executeRuleGroup(andGroup, data);
// Returns true only if age >= 18 AND email is valid AND income >= 25000
```

OR Groups (Any Rule Can Pass)

```
// OR group behavior
RuleGroup orGroup = config.createRuleGroupWithOr("or-group", category, "OR Group", "Any can pass", 20);
// Execution logic:
// - Any rule can evaluate to true for group to pass
```

```
// - If one rule passes, the entire group passes
// - Execution can stop on first success (configurable)
// - Result is true if ANY rule passes

Map<String, Object> data = Map.of("age", 16, "email", null, "income", 30000);
RuleResult result = engine.executeRuleGroup(orGroup, data);
// Returns true if income >= 25000 (even though age and email fail)
```

Integration with Rules Engine

```
// Create rules engine with rule groups
RulesEngineConfiguration config = new RulesEngineConfiguration();

// Add individual rules and rule groups
config.rule("basic-rule").withCondition("#value > 0").build();
config.createRuleGroupWithAnd("validation-group", category, "Validation", "All validations", 10);

RulesEngine engine = new RulesEngine(config);

// Execute all rules and groups for a category
List<RuleResult> results = engine.executeRulesForCategory(category, data);

// Execute specific rule group
RuleGroup group = config.getRuleGroupById("validation-group");
RuleResult groupResult = engine.executeRuleGroup(group, data);
```

Best Practices

Rule Group Organization

- · Use AND groups for validation scenarios where all conditions must be met
- · Use OR groups for eligibility scenarios where any condition is sufficient
- · Group related rules by business domain or functional area
- · Use meaningful IDs and names for easy identification

Performance Optimization

- Enable stop-on-first-failure for AND groups to improve performance
- Enable stop-on-first-success for OR groups when appropriate
- Use parallel-execution for independent rules that can run concurrently
- · Order rules by execution cost (fastest first) for optimal performance

Governance and Maintenance

- Include comprehensive metadata for audit trails
- · Use tags for categorization and filtering
- · Document rule group purpose and business logic
- · Version control rule group configurations

Data Service Configuration

Overview

Data service configuration provides a programmatic way to set up data sources that rules can reference for lookups, enrichments, and data-driven rule evaluation. This approach is particularly useful for testing, demonstrations, and scenarios where you need to configure mock or custom data sources.

DataServiceManager

The DataServiceManager serves as the central orchestration point for all data operations:

```
// Initialize with mock data sources
DataServiceManager dataManager = new DataServiceManager();
dataManager.initializeWithMockData();

// Load custom data sources
dataManager.loadDataSource(new CustomDataSource("ProductsSource", "products"));

// Request data for rule evaluation
List<Product> products = dataManager.requestData("products");
Customer customer = dataManager.requestData("customer");
```

DemoDataServiceManager

For demonstration and testing purposes, the DemoDataServiceManager extends the base manager with pre-configured mock data:

```
public class DemoDataServiceManager extends DataServiceManager {
   @Override
   public DataServiceManager initializeWithMockData() {
        // Create and load mock data sources for various data types
        loadDataSource(new MockDataSource("ProductsDataSource", "products"));
        loadDataSource(new MockDataSource("InventoryDataSource", "inventory"));
        loadDataSource(new MockDataSource("CustomerDataSource", "customer"));
        loadDataSource(new MockDataSource("TemplateCustomerDataSource", "templateCustomer"));
        loadDataSource(new MockDataSource("LookupServicesDataSource", "lookupServices"));
        loadDataSource(new MockDataSource("SourceRecordsDataSource", "sourceRecords"));
        // Add data sources for dynamic matching
        loadDataSource(new MockDataSource("MatchingRecordsDataSource", "matchingRecords"));
        loadDataSource(new MockDataSource("NonMatchingRecordsDataSource", "nonMatchingRecords"));
        return this;
   }
}
```

MockDataSource Implementation

The MockDataSource provides pre-populated test data for various business scenarios:

```
public class MockDataSource implements DataSource {
   private final String name;
   private final String dataType;
   private final Map<String, Object> dataStore = new HashMap<>();

public MockDataSource(String name, String dataType) {
    this.name = name;
    this.dataType = dataType;
```

```
initializeData(); // Populate with test data
}

@Override
public <T> T getData(String dataType, Object... parameters) {
    // Handle special cases like dynamic matching
    if ("matchingRecords".equals(dataType) || "nonMatchingRecords".equals(dataType)) {
        // Dynamic data processing based on parameters
        return processMatchingLogic(parameters);
    }

    return (T) dataStore.get(dataType);
}
```

Integration with Rules

Data services integrate seamlessly with rule evaluation:

```
// Set up data service manager
DemoDataServiceManager dataManager = new DemoDataServiceManager();
dataManager.initializeWithMockData();

// Create rules engine with data context
RulesEngineConfiguration config = new RulesEngineConfiguration();
RulesEngine engine = new RulesEngine(config);

// Get data for rule evaluation
List<Product> products = dataManager.requestData("products");
Customer customer = dataManager.requestData("customer");

// Create evaluation context with data
Map<String, Object> facts = new HashMap<>();
facts.put("products", products);
facts.put("customer", customer);

// Evaluate rules with data context
RuleResult result = engine.evaluate(facts);
```

Custom Data Sources

You can create custom data sources for specific business needs:

```
public class CustomDataSource implements DataSource {
   private final String name;
   private final String dataType;

public CustomDataSource(String name, String dataType) {
        this.name = name;
        this.dataType = dataType;
   }

@Override
public <T> T getData(String dataType, Object... parameters) {
        // Implement custom data retrieval logic
        // Could connect to databases, APIs, files, etc.
        return retrieveCustomData(dataType, parameters);
   }
```

```
@Override
  public boolean supportsDataType(String dataType) {
      return this.dataType.equals(dataType);
  }
}

// Usage
DataServiceManager manager = new DataServiceManager();
manager.loadDataSources(
  new CustomDataSource("CustomProductsSource", "customProducts"),
  new CustomDataSource("CustomCustomerSource", "customCustomer"),
  new CustomDataSource("CustomTradesSource", "customTrades")
);
```

Best Practices

Data Service Organization

- · Use meaningful names for data sources and data types
- · Group related data sources logically
- · Implement proper error handling for data retrieval failures
- · Cache frequently accessed data for performance

Testing and Development

- Use DemoDataServiceManager for demonstrations and testing
- · Create environment-specific data service configurations
- Mock external dependencies during development
- · Validate data integrity before rule evaluation

Production Considerations

- · Replace mock data sources with production implementations
- · Implement proper connection pooling and resource management
- · Add monitoring and health checks for data sources
- · Use appropriate caching strategies for performance

Migration from External Services

Step-by-Step Migration Process

Step 1: Identify Migration Candidates

Analyze your existing lookup services for:

- Small, static datasets (< 100 records)
- · Infrequently changing reference data
- Simple key-value lookups

Step 2: Extract Data

Export data from your existing service:

Step 3: Create YAML Dataset

Convert to YAML format:

```
enrichments:
- id: "currency-enrichment"
type: "lookup-enrichment"
lookup-config:
lookup-dataset:
type: "inline"
key-field: "code"
data:
- code: "USD"
name: "US Dollar"
region: "North America"
- code: "EUR"
name: "Euro"
region: "Europe"
```

Step 4: Update Configuration

Replace service calls with enrichment:

```
// After: YAML dataset enrichment
RulesEngineConfiguration config = YamlConfigurationLoader.load("config.yaml");
RulesEngine engine = new RulesEngine(config);
RuleResult result = engine.evaluate(data); // Enrichment happens automatically
```

Step 5: Test and Validate

Ensure the migration works correctly:

```
@Test
public void testCurrencyEnrichment() {
    Map<String, Object> data = Map.of("currency", "USD");
    RuleResult result = engine.evaluate(data);

    assertEquals("US Dollar", result.getEnrichedData().get("currencyName"));
    assertEquals("North America", result.getEnrichedData().get("currencyRegion"));
}
```

Best Practices

Following these best practices will help you build maintainable, performant, and reliable rule-based systems with APEX. These recommendations come from real-world experience and will save you time and effort in the long run.

Configuration Organization

Good organization makes your rules easier to understand, maintain, and debug.

File Organization:

- Use external dataset files for reusable data: If multiple configurations need the same lookup data, put it in a separate file
 that can be shared
- · Keep inline datasets small: Limit inline datasets to less than 50 records to keep configuration files readable
- Use meaningful IDs and names: Choose descriptive identifiers like "customer-age-validation" instead of "rule1"
- Include metadata for documentation: Add owner, purpose, and creation date information to help future maintainers

Naming Conventions:

- Use consistent naming patterns across your organization
- Include the business domain in rule IDs (e.g., "finance-trade-validation", "customer-eligibility-check")
- · Use descriptive messages that help users understand what went wrong

Performance Optimization

APEX is designed for high performance, but following these practices will ensure optimal speed.

Caching Strategy:

- Enable caching for frequently accessed datasets: Turn on caching for lookup data that's used often
- Use appropriate cache TTL values: Set cache expiration times based on how often your data changes
 - o Static data (countries, currencies): 24 hours or more
 - Semi-static data (product categories): 1-4 hours
 - o Dynamic data: 5-30 minutes
- Monitor performance metrics: Use APEX's built-in monitoring to identify slow rules or enrichments
- Preload datasets when possible: Load reference data at startup rather than on first use

Rule Optimization:

- · Order rules by execution cost (fastest first) when using rule groups
- · Use specific conditions to avoid unnecessary rule evaluations
- · Consider using rule chains for complex multi-step logic

Maintenance and Governance

Proper maintenance practices prevent technical debt and ensure long-term success.

Version Control:

- Version control all configuration files: Treat YAML configurations like code use Git or similar
- Use environment-specific configurations: Have separate configurations for development, testing, and production
- Document dataset sources and update procedures: Record where data comes from and how to update it
- Regular review and cleanup of unused datasets: Remove obsolete rules and datasets to keep configurations clean

Change Management:

- · Test configuration changes in non-production environments first
- Use meaningful commit messages when updating configurations
- Consider the impact of rule changes on existing processes
- · Maintain a changelog for significant rule modifications

Advanced Rule Patterns

For complex business scenarios requiring rule dependencies and chaining, the Rules Engine supports sophisticated patterns where rules depend on results of previous rules. These patterns are essential for multi-stage workflows and decision trees.

Available Patterns

- 1. Conditional Chaining Execute Rule B only if Rule A triggers
- 2. Sequential Dependency Each rule builds upon results from the previous rule
- 3. Result-Based Routing Route to different rule sets based on previous results
- 4. Accumulative Chaining Build up a score/result across multiple rules
- 5. Complex Financial Workflow Real-world nested rule scenarios with multi-stage processing
- 6. Fluent Rule Builder Compose rules with conditional execution paths using fluent API

Quick Example: Conditional Chaining

```
// Rule A: Check if customer qualifies for high-value processing
Rule ruleA = new Rule(
   "HighValueCustomerCheck",
   "#customerType == 'PREMIUM' && #transactionAmount > 100000",
    "Customer qualifies for high-value processing"
);
// Execute Rule A first
List<RuleResult> resultsA = ruleEngineService.evaluateRules(
   Arrays.asList(ruleA), createEvaluationContext(context));
// Conditional execution of Rule B based on Rule A result
if (resultsA.get(0).isTriggered()) {
   Rule ruleB = new Rule(
        "EnhancedDueDiligenceCheck",
        "#accountAge >= 3",
        "Enhanced due diligence check passed"
   );
   // Execute enhanced validation only when needed
}
```

Quick Example: Accumulative Chaining

```
rule-chains:
    - id: "credit-scoring"
    pattern: "accumulative-chaining"
    configuration:
        accumulator-variable: "totalScore"
        initial-value: 0
        accumulation-rules:
        - id: "credit-score-component"
        condition: "#creditScore >= 700 ? 25 : (#creditScore >= 650 ? 15 : 10)"
        message: "Credit score component"
```

```
weight: 1.0
- id: "income-component"
  condition: "#annualIncome >= 80000 ? 20 : 15"
  message: "Income component"
  weight: 1.0
final-decision-rule:
  id: "loan-decision"
  condition: "#totalScore >= 60 ? 'APPROVED' : 'DENIED'"
  message: "Final loan decision"
```

When to Use Advanced Patterns

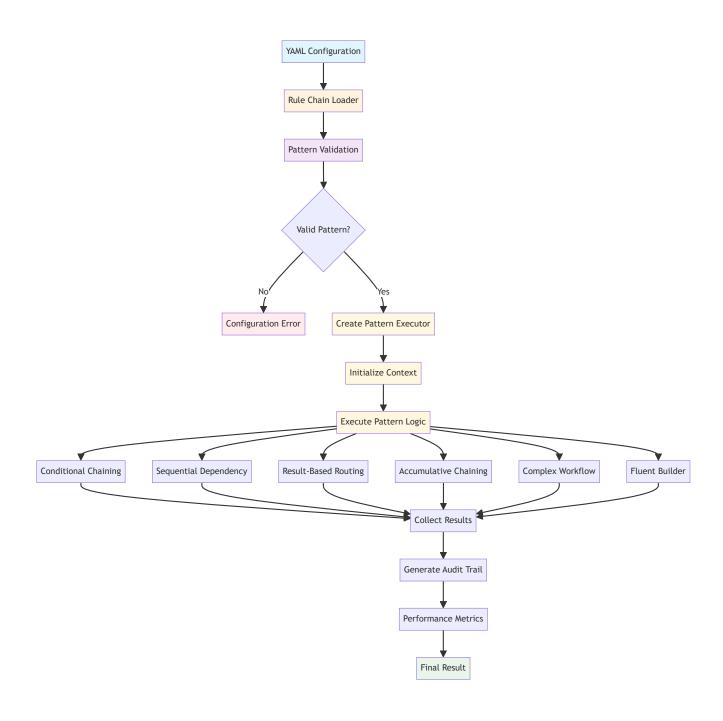
Use advanced rule patterns when you need:

- Multi-stage approval processes
- · Risk-based processing with different validation paths
- · Complex financial workflows with dependencies
- · Decision trees with conditional branching
- Accumulative scoring systems
- · Performance optimization through conditional execution

Documentation Reference

Rule Chains Configuration

Rule chains are configured in the rule-chains section of your YAML configuration file. Each rule chain specifies a pattern and its configuration:



```
rule-chains:
    - id: "my-rule-chain"
    name: "My Rule Chain"
    description: "Description of what this chain does"
    pattern: "conditional-chaining" # One of the 6 supported patterns
    enabled: true
    priority: 10
    category: "business-logic"
    configuration:
        # Pattern-specific configuration goes here
```

Pattern-Specific Configuration Examples

Execute expensive or specialized rules only when certain conditions are met:

```
rule-chains:
    - id: "high-value-processing"
    pattern: "conditional-chaining"
    configuration:
        trigger-rule:
        condition: "#customerType == 'PREMIUM' && #transactionAmount > 100000"
        message: "High-value customer transaction detected"
        conditional-rules:
        on-trigger:
        - condition: "#accountAge >= 3"
            message: "Enhanced due diligence check"
        on-no-trigger:
        - condition: "true"
        message: "Standard processing applied"
```

Pattern 2: Sequential Dependency

Build processing pipelines where each stage uses results from previous stages:

```
rule-chains:
 - id: "discount-pipeline"
   pattern: "sequential-dependency"
   configuration:
      stages:
        - stage: 1
         name: "Base Discount"
         rule:
            condition: "#customerTier == 'GOLD' ? 0.15 : 0.05"
           message: "Base discount calculated"
         output-variable: "baseDiscount"
        - stage: 2
          name: "Regional Multiplier"
          rule:
            condition: "#region == 'US' ? #baseDiscount * 1.2 : #baseDiscount"
            message: "Regional multiplier applied"
          output-variable: "finalDiscount"
```

Pattern 3: Result-Based Routing

Route to different rule sets based on intermediate results:

Pattern 4: Accumulative Chaining with Weight-Based Rule Selection

Build up scores across multiple criteria with weighted components and intelligent rule selection:

```
rule-chains:
  - id: "advanced-credit-scoring"
   pattern: "accumulative-chaining"
   configuration:
     accumulator-variable: "totalScore"
     initial-value: 0
     # NEW: Rule Selection Strategy
     rule-selection:
        strategy: "weight-threshold" # Execute only high-importance rules
        weight-threshold: 0.7
      accumulation-rules:
        - id: "credit-score-component"
          condition: "#creditScore >= 700 ? 25 : (#creditScore >= 650 ? 15 : 10)"
         message: "Credit score component"
         weight: 0.9
                          # High importance - will be executed
         priority: "HIGH"
        - id: "income-component"
          condition: "#annualIncome >= 80000 ? 20 : (#annualIncome >= 60000 ? 15 : 10)"
          message: "Income component"
         weight: 0.8
                         # High importance - will be executed
          priority: "HIGH"
        - id: "employment-component"
          condition: "#employmentYears >= 5 ? 15 : (#employmentYears >= 2 ? 10 : 5)"
         message: "Employment component"
         weight: 0.6
                          # Below threshold - will be skipped
         priority: "MEDIUM"
        - id: "debt-ratio-component"
          condition: "(#existingDebt / #annualIncome) < 0.2 ? 10 : 0"</pre>
          message: "Debt-to-income ratio component"
          weight: 0.5
                         # Below threshold - will be skipped
          priority: "LOW"
      final-decision-rule:
        condition: "#totalScore >= 40 ? 'APPROVED' : (#totalScore >= 25 ? 'CONDITIONAL' : 'DENIED')"
        message: "Final loan decision"
```

Rule Selection Strategies:

1. Weight Threshold: Execute only rules above a weight threshold

```
rule-selection:
   strategy: "weight-threshold"
   weight-threshold: 0.7 # Only rules with weight >= 0.7
```

2. Top Weighted: Execute the N highest-weighted rules

```
rule-selection:
   strategy: "top-weighted"
   max-rules: 3  # Execute top 3 rules by weight
```

3. Priority Based: Execute rules based on priority levels

```
rule-selection:
   strategy: "priority-based"
   min-priority: "MEDIUM" # Execute HIGH and MEDIUM priority rules
```

4. Dynamic Threshold: Calculate threshold based on context

```
rule-selection:
   strategy: "dynamic-threshold"
   threshold-expression: "#riskLevel == 'HIGH' ? 0.8 : 0.6"
```

Advanced Accumulative Features:

1. Weighted Components: Different importance levels for components

```
- id: "critical-component"
  condition: "#value > 100 ? 20 : 10"
  weight: 2.0 # This component has double impact
  priority: "HIGH"
```

2. Negative Scoring: Penalties that reduce the total score

```
- id: "risk-penalty"
  condition: "#riskFactors > 3 ? -15 : 0"  # Subtract points for high risk
  weight: 1.0
  priority: "MEDIUM"
```

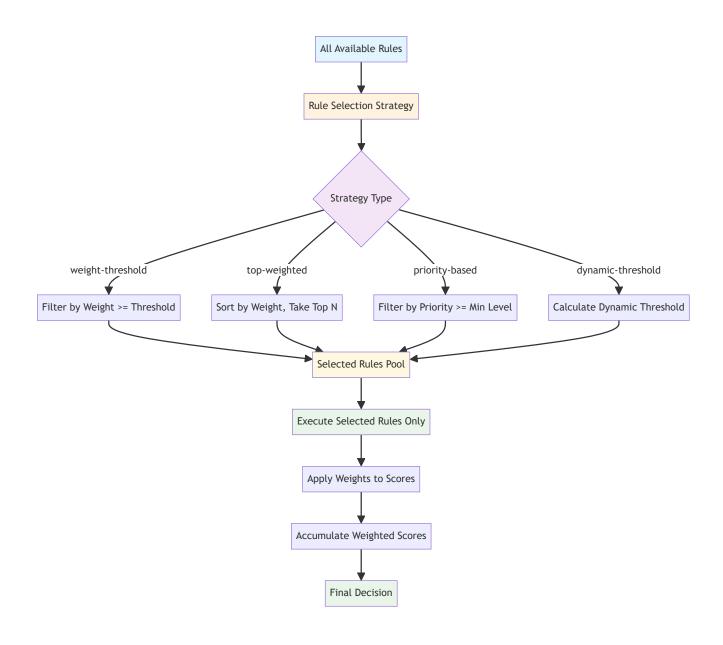
3. Flexible Initial Values: Start with a base score

```
configuration:
  initial-value: 50 # Start with 50 points instead of 0
```

4. Complex Conditional Logic: Multi-tier scoring within rules

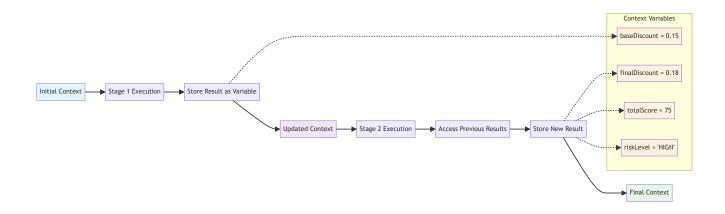
```
- id: "tiered-scoring"
condition: "#value >= 100 ? 30 : (#value >= 75 ? 20 : (#value >= 50 ? 10 : 0))"
weight: 0.9
priority: "HIGH"
```

Weight-Based Rule Selection Flow:



Execution Context and Variable Propagation

Rule chains maintain an execution context that allows results from one stage to be used in subsequent stages:



• Stage Results: Each stage can store results using output-variable

- · Accumulator Variables: Accumulative chains maintain running totals
- · Context Variables: All results are available as SpEL variables in later rules

Example of variable usage:

```
# Stage 1 stores result as 'baseDiscount'
- stage: 1
rule:
    condition: "#customerTier == 'GOLD' ? 0.15 : 0.05"
    output-variable: "baseDiscount"

# Stage 2 uses 'baseDiscount' from Stage 1
- stage: 2
rule:
    condition: "#baseDiscount * 1.2" # Uses result from Stage 1
    output-variable: "finalDiscount"
```

Execution Results and Monitoring

Rule chains provide comprehensive execution results and monitoring capabilities:

Accessing Execution Results

```
// Execute a rule chain
RuleChainResult result = ruleChainExecutor.executeRuleChain(ruleChain, context);
// Check overall success
if (result.isSuccessful()) {
   System.out.println("Chain executed successfully: " + result.getFinalOutcome());
} else {
   System.out.println("Chain failed: " + result.getErrorMessage());
// Access execution path
List<String> executionPath = result.getExecutionPath();
System.out.println("Rules executed: " + String.join(" → ", executionPath));
// Access stage results (for sequential dependency and accumulative chaining)
Object stageResult = result.getStageResult("baseDiscount");
System.out.println("Base discount calculated: " + stageResult);
// Access performance metrics
long executionTime = result.getExecutionTimeMillis();
int rulesExecuted = result.getExecutedRulesCount();
int rulesTriggered = result.getTriggeredRulesCount();
```

Pattern-Specific Result Access

Sequential Dependency Results:

```
// Access results from each stage
Double baseDiscount = (Double) result.getStageResult("baseDiscount");
Double finalDiscount = (Double) result.getStageResult("finalDiscount");
BigDecimal finalAmount = (BigDecimal) result.getStageResult("finalAmount");
```

Accumulative Chaining Results:

```
// Access component scores
Double creditComponent = (Double) result.getStageResult("component_1_credit-score-component_score");
Double incomeComponent = (Double) result.getStageResult("component_2_income-component_score");

// Access weighted scores
Double creditWeighted = (Double) result.getStageResult("component_1_credit-score-component_weighted");

// Access accumulator progression
Double initialScore = (Double) result.getStageResult("totalScore_initial");
Double finalScore = (Double) result.getStageResult("totalScore_final");
```

Result-Based Routing Results:

```
// Access routing decision
String routeKey = (String) result.getStageResult("routeKey");
String routeExecutionResult = (String) result.getStageResult("routeExecutionResult");
Integer routeExecutedRules = (Integer) result.getStageResult("routeExecutedRules");
```

Error Handling and Validation

Rule chains include comprehensive validation and error handling:

- Configuration Validation: Pattern-specific validation ensures required fields are present
- . Runtime Error Handling: Individual rule failures don't stop the entire chain
- Execution Tracking: Full audit trail of which rules executed and their results
- Performance Monitoring: Execution time tracking for optimization
- Detailed Error Messages: Specific error information for troubleshooting
- Graceful Degradation: Chains continue executing even when individual rules fail

Practical Examples

Credit Application Processing

Combine multiple patterns for comprehensive credit application processing:

```
rule-chains:
 # First, determine application complexity
  - id: "application-complexity-routing"
   pattern: "result-based-routing"
   configuration:
       condition: "#loanAmount > 500000 || #applicantType == 'BUSINESS' ? 'COMPLEX' : 'SIMPLE'"
     routes:
       COMPLEX:
         rules:
            - condition: "#documentationComplete == true"
             message: "Complex application documentation check"
        STMPLE.
          rules:
            - condition: "#basicInfoComplete == true"
              message: "Simple application basic info check"
 # Then, calculate credit score
  - id: "credit-score-calculation"
   pattern: "accumulative-chaining"
   configuration:
```

```
accumulator-variable: "creditScore"
    initial-value: 0
    accumulation-rules:
      - condition: "#creditHistory >= 700 ? 30 : 15"
        weight: 1.0
      - condition: "#incomeStability >= 0.8 ? 25 : 10"
        weight: 1.0
      - condition: "#debtToIncomeRatio < 0.3 ? 20 : 0"</pre>
        weight: 1.0
    final-decision-rule:
      condition: "#creditScore >= 60 ? 'APPROVED' : 'DENIED'"
# Finally, apply conditional enhanced processing for high-value loans
- id: "enhanced-processing"
  pattern: "conditional-chaining"
  configuration:
    trigger-rule:
      condition: "#loanAmount > 1000000 && #creditScore >= 60"
    conditional-rules:
      on-trigger:
        - condition: "#manualReviewRequired = true"
          message: "High-value loan requires manual review"
```

Performance-Based Pricing

Use sequential dependency for complex pricing calculations:

```
rule-chains:
  - id: "performance-pricing"
   pattern: "sequential-dependency"
   configuration:
     stages:
        - stage: 1
         name: "Base Rate Calculation"
            condition: "#customerTier == 'PLATINUM' ? 0.02 : (#customerTier == 'GOLD' ? 0.025 : 0.03)"
          output-variable: "baseRate"
        - stage: 2
         name: "Volume Discount"
          rule:
            condition: "#tradingVolume > 10000000 ? #baseRate * 0.8 : (#tradingVolume > 5000000 ? #baseRate * 0.9 : #base
          output-variable: "discountedRate"
        - stage: 3
          name: "Relationship Bonus"
            condition: "#relationshipYears > 5 ? #discountedRate * 0.95 : #discountedRate"
          output-variable: "finalRate"
```

Comprehensive Financial Services Example

Combining all 4 implemented patterns for complete loan processing:

```
rule-chains:
    # Step 1: Route based on loan type and amount
    - id: "loan-application-routing"
    pattern: "result-based-routing"
    priority: 10
    configuration:
        router-rule:
            condition: "#loanType == 'MORTGAGE' && #loanAmount > 500000 ? 'COMPLEX_MORTGAGE' : (#loanType == 'PERSONAL' ? 'PE
```

```
output-variable: "loanCategory"
    routes:
      COMPLEX_MORTGAGE:
          - condition: "#propertyAppraisal != null && #incomeVerification == true"
            message: "Complex mortgage documentation verified"
      PERSONAL_LOAN:
        rules:
          - condition: "#loanAmount <= 50000"</pre>
            message: "Personal loan within limits"
      STANDARD LOAN:
        rules:
          - condition: "#basicDocumentation == true"
            message: "Standard loan documentation complete"
# Step 2: Calculate comprehensive credit score
- id: "comprehensive-credit-scoring"
 pattern: "accumulative-chaining"
 priority: 20
 configuration:
   accumulator-variable: "creditScore"
   initial-value: 0
   accumulation-rules:
      - id: "credit-history"
        condition: "#creditRating >= 750 ? 35 : (#creditRating >= 700 ? 30 : (#creditRating >= 650 ? 20 : 10))"
        weight: 1.5 # Credit history is most important
      - id: "income-stability"
        condition: "#annualIncome >= 100000 ? 25 : (#annualIncome >= 75000 ? 20 : 15)"
       weight: 1.2
      - id: "employment-history"
        condition: "#employmentYears >= 5 ? 20 : (#employmentYears >= 2 ? 15 : 10)"
        weight: 1.0
      - id: "debt-burden"
       condition: "(#totalDebt / #annualIncome) < 0.2 ? 15 : ((#totalDebt / #annualIncome) < 0.4 ? 5 : -10)"</pre>
       weight: 1.0
      - id: "assets"
        condition: "#liquidAssets >= 50000 ? 10 : (#liquidAssets >= 25000 ? 5 : 0)"
        weight: 0.8
   final-decision-rule:
      condition: "#creditScore >= 80 ? 'EXCELLENT' : (#creditScore >= 60 ? 'GODD' : (#creditScore >= 40 ? 'FAIR' : 'POO
# Step 3: Apply conditional enhanced processing for high-value loans
- id: "enhanced-processing"
 pattern: "conditional-chaining"
 priority: 30
 configuration:
   trigger-rule:
      condition: "#loanAmount > 1000000 && #creditScore >= 60"
      message: "High-value loan with acceptable credit"
   conditional-rules:
     on-trigger:
       - condition: "#manualUnderwritingRequired = true"
          message: "Manual underwriting required for high-value loan"
        - condition: "#seniorApprovalRequired = true"
         message: "Senior management approval required"
      on-no-trigger:
        - condition: "#standardProcessing = true"
          message: "Standard automated processing"
# Step 4: Calculate final terms based on all previous results
- id: "loan-terms-calculation"
 pattern: "sequential-dependency"
 priority: 40
 configuration:
   stages:
      - stage: 1
```

```
name: "Base Interest Rate"
rule:
    condition: "#creditScore >= 80 ? 0.035 : (#creditScore >= 60 ? 0.045 : (#creditScore >= 40 ? 0.055 : 0.065))"
    output-variable: "baseRate"
- stage: 2
    name: "Loan Amount Adjustment"
    rule:
        condition: "#loanAmount > 1000000 ? #baseRate + 0.005 : #baseRate"
        output-variable: "adjustedRate"
- stage: 3
    name: "Final Rate with Fees"
    rule:
        condition: "#adjustedRate + (#loanAmount * 0.001 / #loanAmount)"
        output-variable: "finalRate"
```

Complex Financial Workflow Example

Pattern 5 for multi-stage processing with dependencies:

```
rule-chains:
  - id: "trade-settlement-workflow"
   pattern: "complex-workflow"
   configuration:
     stages:
        - stage: "trade-validation"
          name: "Trade Data Validation"
            - condition: "#tradeType != null && #notionalAmount != null && #counterparty != null"
              message: "Basic trade data validation"
          failure-action: "terminate"
        - stage: "risk-assessment"
          name: "Risk Assessment"
          depends-on: ["trade-validation"]
          rules:
            - condition: "#notionalAmount > 1000000 && #marketVolatility > 0.2 ? 'HIGH' : 'MEDIUM'"
              message: "Risk level assessment"
          output-variable: "riskLevel"
        - stage: "approval-workflow"
          name: "Approval Workflow"
          depends-on: ["risk-assessment"]
          conditional-execution:
            condition: "#riskLevel == 'HIGH'"
            on-true:
                - condition: "#seniorApprovalObtained == true"
                  message: "Senior approval required for high-risk trades"
            on-false:
                - condition: "true"
                 message: "Standard approval applied"
        - stage: "settlement-calculation"
          name: "Settlement Processing"
          depends-on: ["approval-workflow"]
            - condition: "#tradeType == 'DERIVATIVE' ? 5 : (#tradeType == 'EQUITY' ? 3 : 2)"
              message: "Settlement days calculated"
          output-variable: "settlementDays"
```

Fluent Decision Tree Example

Pattern 6 for complex decision trees:

```
rule-chains:
  - id: "customer-onboarding-tree"
   pattern: "fluent-builder"
   configuration:
     root-rule:
        id: "customer-tier-check"
        condition: "#customerType == 'VIP' || #customerType == 'PREMIUM'"
       message: "High-tier customer detected"
       on-success:
          rule:
            id: "account-value-check"
            condition: "#initialDeposit > 100000"
            message: "High-value account detected"
            on-success:
              rule:
                id: "expedited-onboarding"
                condition: "true"
                message: "Expedited onboarding approved"
            on-failure:
              rule:
                id: "standard-premium-onboarding"
                condition: "true"
                message: "Standard premium onboarding"
       on-failure:
          rule:
            id: "standard-customer-check"
            condition: "#initialDeposit > 1000"
            message: "Standard customer validation"
            on-success:
              rule:
                id: "standard-onboarding"
                condition: "true"
                message: "Standard onboarding approved"
            on-failure:
              rule:
                id: "minimum-deposit-required"
                condition: "true"
                message: "Minimum deposit requirement not met"
```

This comprehensive example demonstrates how all 6 implemented patterns work together to create sophisticated business processing systems with routing, scoring, conditional processing, sequential calculations, complex workflows, and decision trees.

Integration with Existing Rules

Rule chains work alongside traditional rules and rule groups:

```
# Traditional rules continue to work
rules:
    - id: "basic-rule"
        condition: "#amount > 0"
        message: "Basic validation"

# Rule groups continue to work
rule-groups:
    - id: "validation-group"
        rule-ids: ["basic-rule"]

# Rule chains add advanced patterns
rule-chains:
    - id: "advanced-processing"
        pattern: "sequential-dependency"
```

For complete implementation details, examples, and architecture information, see the **Technical Reference Guide** section on "Nested Rules and Rule Chaining Patterns".

Configuration Standards and Validation

Overview

This section provides comprehensive standards for YAML configuration management in APEX, including mandatory metadata requirements, file type specifications, scenario management, and validation procedures. These standards ensure consistency, maintainability, and regulatory compliance across all APEX configurations.

File Type System

APEX uses a standardized file type system to categorize and validate different kinds of YAML configurations:

Туре	Purpose	Additional Required Fields	Content Validation
scenario	Data type routing	business-domain , owner	scenario section with data-types and rule- configurations
scenario- registry	Central registry	created-by	scenario-registry list with valid entries
bootstrap	Complete demos	business-domain , created-by	rule-chains Or categories sections
rule-config	Reusable rules	author	rules , enrichments , or rule-chains sections
dataset	Reference data	source	data , countries , or dataset sections
enrichment	Data enrichment	author	Enrichment-specific content
rule-chain	Sequential rules	author	Rule chain definitions

Complete Metadata Examples by Type

```
Bootstrap Files ( type: "bootstrap" )
```

```
metadata:
   name: "OTC Options Bootstrap Configuration"
   version: "1.0.0"
   description: "Complete OTC Options processing demonstration"
   type: "bootstrap"
   business-domain: "Derivatives Trading"  # Required: Business context
```

```
created-by: "bootstrap.admin@company.com" # Required: Creator identification
     created: "2025-08-02"
                                              # Optional: Creation date
     tags: ["derivatives", "options", "demo"] # Optional: Classification tags
Rule Configuration Files ( type: "rule-config" )
   metadata:
    name: "Financial Validation Rules"
     version: "1.0.0"
     description: "Comprehensive validation rules for financial instruments"
     type: "rule-config"
     author: "rules.team@company.com"
                                               # Required: Rule authorship
     business-domain: "Financial Services" # Optional: Business context
    created: "2025-08-02" # Optional: Creation date
last-modified: "2025-08-02" # Optional: Last modification
Dataset Files ( type: "dataset" )
   metadata:
    name: "Countries Lookup Dataset"
     version: "1.0.0"
     description: "Country codes with currency and timezone data"
     type: "dataset"
     source: "ISO 3166-1 alpha-2 country codes" # Required: Data source
    last-updated: "2025-08-02"  # Optional: Data freshness data-classification: "Public"  # Optional: Data sensitivity
Scenario Registry ( type: "scenario-registry" )
   metadata:
     name: "Scenario Registry Configuration"
     version: "1.0.0"
     description: "Central registry of all available scenarios"
     type: "scenario-registry"
     created-by: "registry.admin@company.com" # Required: Registry manager
    created: "2025-08-02" # Optional: Creation data-updated: "2025-08-02" # Optional: Last update
                                               # Optional: Creation date
```

Financial Services Compliance Standards

For financial services environments, additional metadata fields support regulatory compliance and audit requirements:

```
metadata:
 name: "EMIR Reporting Validation Rules"
 version: "1.0.0"
 description: "Validation rules for EMIR regulatory reporting"
 type: "rule-config"
 author: "regulatory.team@firm.com"
 business-domain: "Regulatory Reporting"
 regulatory-scope: "European Union (EMIR)"
                                               # Required: Regulatory context
  compliance-reviewed: true
                                               # Required: Compliance approval
  compliance-reviewer: "compliance@firm.com" # Required: Who reviewed
                                # Required: When reviewed
 compliance-date: "2025-08-02"
 risk-approved: true
                                             # Required: Risk approval
 risk-reviewer: "risk@firm.com"
                                           # Required: Risk approver
```

Regulatory Compliance Fields:

- regulatory-scope: Applicable regulations and jurisdictions
- compliance-reviewed: Boolean indicating compliance team approval
- compliance-reviewer : Email of compliance reviewer
- compliance-date : Date of compliance review
- risk-approved: Boolean indicating risk team approval
- risk-reviewer: Email of risk approver
- operational-impact : Impact level (Low, Medium, High, Critical)

Configuration Standards and Conventions

1. Naming Conventions

File Naming Patterns:

- Scenarios: {domain}-{type}-{variant}-scenario.yaml
 - Example: derivatives-otc-options-standard-scenario.yaml
- Bootstrap: {domain}-{use-case}-bootstrap.yaml
 - o Example: derivatives-otc-options-bootstrap.yaml
- Rule Config: {domain}-{purpose}-rules.yaml
 - Example: financial-validation-rules.yaml

Scenario IDs:

- Use kebab-case: otc-options-standard
- Include business domain: derivatives-otc-options-standard
- Environment suffix if needed: otc-options-dev , otc-options-prod

2. Version Management

- Use semantic versioning (1.0.0, 1.1.0, 2.0.0)
- · Increment major version for breaking changes
- Increment minor version for new features
- Increment patch version for bug fixes

3. Documentation Standards

- · Provide clear, descriptive names that explain the file's purpose
- Include comprehensive descriptions for better documentation
- Document business context and ownership for maintainability
- · Add tags for categorization to improve discoverability

4. Validation Integration

- Run validation as part of CI/CD pipeline to catch errors early
- · Validate all files before deployment to ensure quality
- Generate validation reports for compliance documentation
- Set up automated alerts for validation failures for quick response

Implementation Checklists

For New YAML Files

When creating new YAML configuration files, follow this checklist: Include all required metadata fields (name , version , description , type) Use correct type value for file purpose (scenario, bootstrap, rule-config, etc.) Add type-specific required fields (author, business-domain, owner, etc.) Follow naming conventions for files and IDs Include business domain and ownership information Add regulatory scope if applicable for financial services ■ Validate file before committing using YAML validation tools For Existing Files When updating existing YAML files to meet current standards: Add missing type field to metadata section Verify all required fields are present for the file type Update to use standardized type values (e.g., rule-config not rules) Add missing business context fields (business-domain, owner, etc.) Run comprehensive validation to check for errors Fix any validation errors before deployment For CI/CD Integration To integrate YAML validation into your development workflow: Add YAML validation to build pipeline using YamlMetadataValidator Configure validation failure alerts to notify teams of issues ■ Generate validation reports for compliance documentation Set up automated quality checks for pull requests Document validation procedures for team members Validation Error Examples and Solutions **Common Validation Errors:** ERROR: Missing required metadata field: type SOLUTION: Add type field to metadata section with appropriate value

ERROR: Invalid file type: invalid-type
SOLUTION: Use one of: scenario, scenario-registry, bootstrap, rule-config, dataset, enrichment, rule-chain
ERROR: Scenario files must have a 'scenario' section
SOLUTION: Add scenario section with scenario-id, data-types, and rule-configurations

ERROR: Missing required field for type 'scenario': business-domain

SOLUTION: Add business-domain field to scenario metadata

Automated Validation Setup

Basic Validation in Java:

```
// Validate a single file
YamlMetadataValidator validator = new YamlMetadataValidator();
YamlValidationResult result = validator.validateFile("scenarios/otc-options-scenario.yaml");

if (result.isValid()) {
    System.out.println(" / File is valid");
} else {
    System.out.println(" / Validation errors:");
    result.getErrors().forEach(error -> System.out.println(" - " + error));
}

// Validate multiple files
List<String> files = List.of("scenarios/scenario1.yaml", "config/rules.yaml");
YamlValidationSummary summary = validator.validateFiles(files);
String report = summary.getReport(); // Comprehensive validation report
```

CI/CD Pipeline Integration:

```
# .github/workflows/yaml-validation.yml
name: YAML Configuration Validation
on:
 push:
   paths:
     - '**/*.yaml'
     - '**/*.yml'
 pull_request:
   paths:
     - '**/*.yaml'
      - '**/*.yml'
jobs:
 validate-yaml:
   runs-on: ubuntu-latest
     - uses: actions/checkout@v3
     - name: Set up Java
       uses: actions/setup-java@v3
          java-version: '17'
         distribution: 'temurin'
     - name: Run YAML Validation
       run:
         cd apex-core
         mvn test -Dtest=ComprehensiveYamlValidationTest
```

Quality Assurance Guidelines

Pre-Deployment Validation

Before deploying any YAML configuration changes:

- 1. Run comprehensive validation on all affected files
- 2. Check dependency chains for missing references
- 3. Verify business logic with domain experts
- 4. Test in lower environments before production
- 5. Document changes in version control

Ongoing Maintenance

For long-term configuration quality:

- 1. Regular validation audits to catch drift
- 2. Automated quality reports for management visibility
- 3. Team training on configuration standards
- 4. Documentation updates as standards evolve
- 5. Compliance reviews for regulatory requirements

Getting Help

Common Issues

Configuration Issues:

- Configuration not loading: Check YAML syntax and file paths
- Missing metadata errors: Ensure all required metadata fields are present (see Configuration Standards)
- · Validation failures: Use the implementation checklists to verify file structure
- Type errors: Verify correct type value for your file purpose

Runtime Issues:

- · Enrichment not working: Verify condition expressions and field mappings
- Performance issues: Enable caching and monitor metrics
- · Data not found: Check key field matching and default values
- · Scenario not found: Verify scenario registry and data type mappings

Documentation Resources

Core Guides:

- Configuration Standards and Validation: Comprehensive standards for YAML files
- Technical Implementation Guide: Architecture details and advanced patterns
- Financial Services Guide: Domain-specific examples and compliance requirements

Quick References:

- Mandatory Metadata Attributes: Required fields for all file types
- Implementation Checklists: Step-by-step validation guides
- Configuration examples and templates: Working examples for common use cases

Support

Self-Service:

- Check validation errors using YamlMetadataValidator for specific error messages
- Review configuration standards for proper file structure and metadata
- Use implementation checklists to verify your configurations meet requirements

Community Support:

- · Create GitHub issues for bugs or feature requests
- · Check existing documentation for common solutions
- Review configuration examples for similar use cases

Enterprise Support:

- Configuration validation services for large-scale deployments
- · Custom training on APEX configuration standards
- Regulatory compliance consulting for financial services